

# Support Vector Machines and Kernel methods

Morteza H. Chehreghani

`morteza.chehreghani@chalmers.se`

Department of Computer Science and Engineering  
Chalmers University

February 2, 2024

# Reference

The content and the slides are adapted from

S. Rogers and M. Girolami, A First Course in Machine Learning (FCML), 2nd edition, Chapman & Hall/CRC 2016, ISBN: 9781498738484

# Classification syllabus

- ▶ 4 classification algorithms.
- ▶ Of which:
  - ▶ 2 are probabilistic.
    - ▶ Bayes classifier.
    - ▶ Logistic regression.
  - ▶ 2 non-probabilistic.
    - ▶ K-nearest neighbours.
    - ▶ Support Vector Machines (SVM).
- ▶ There are many others!

# Topics ...

- ▶ Linear SVM
- ▶ Soft-Margin SVM
- ▶ Kernels - Kernel SVM
- ▶ Classifier Performance

# Topics ...

- ▶ **Linear SVM**
- ▶ Soft-Margin SVM
- ▶ Kernels - Kernel SVM
- ▶ Classifier Performance

# The margin


- ▶ We have seen several algorithms where we find the parameters that optimise something:
  - ▶ Minimise the loss.
  - ▶ Maximise the likelihood.
  - ▶ Maximise the posterior (MAP).

# The margin

- ▶ We have seen several algorithms where we find the parameters that optimise something:
  - ▶ Minimise the loss.
  - ▶ Maximise the likelihood.
  - ▶ Maximise the posterior (MAP).
- ▶ The Support Vector Machine (SVM) is no different:
- ▶ It finds the *decision boundary* that maximises the **margin**.

# Some data

- We'll 'think' in 2-dimensions.




SVM is a binary classifier.  
 $N$  data points, each with  
attributes  $\mathbf{x} = [x_1, x_2]^T$  and  
target  $t = \pm 1$



# Some data

- ▶ We'll 'think' in 2-dimensions.



SVM is a binary classifier.  
 $N$  data points, each with  
attributes  $\mathbf{x} = [x_1, x_2]^T$  and  
target  $t = \pm 1$

- ▶ A linear *decision boundary* can be represented as a straight line:


$$\mathbf{w}^T \mathbf{x} + b = 0$$

↓  
2-dim

$$\mathbf{w}^T \mathbf{x} = w_1 x_1 + w_2 x_2$$

## Some data

- ▶ We'll 'think' in 2-dimensions.



SVM is a binary classifier.  
 $N$  data points, each with  
attributes  $\mathbf{x} = [x_1, x_2]^T$  and  
target  $t = \pm 1$


- ▶ A linear *decision boundary* can be represented as a straight line:

$$\mathbf{w}^T \mathbf{x} + b = 0$$

- ▶ Our task is to find  $\mathbf{w}$  and  $b$

# Some data

- ▶ We'll 'think' in 2-dimensions.



SVM is a binary classifier.  
 $N$  data points, each with  
attributes  $\mathbf{x} = [x_1, x_2]^T$  and  
target  $t = \pm 1$

- ▶ A linear *decision boundary* can be represented as a straight line:

$$\mathbf{w}^T \mathbf{x} + b = 0$$

- ▶ Our task is to find  $\mathbf{w}$  and  $b$
- ▶ Once we have these, classification is easy:

$$\left\{ \begin{array}{ll} \mathbf{w}^T \mathbf{x}_{\text{new}} + b > 0 & : t_{\text{new}} = 1 \\ \mathbf{w}^T \mathbf{x}_{\text{new}} + b < 0 & : t_{\text{new}} = -1 \end{array} \right.$$


- ▶ i.e.  $t_{\text{new}} = \text{sign}(\mathbf{w}^T \mathbf{x}_{\text{new}} + b)$

# The margin

- ▶ How do we choose  $\mathbf{w}$  and  $b$ ?
- ▶ Need a quantity to optimise!


# The margin

- ▶ How do we choose  $\mathbf{w}$  and  $b$ ?
- ▶ Need a quantity to optimise!
- ▶ Use the **margin**,  $\gamma$
- ▶ Maximise it!




# The margin

- ▶ How do we choose  $\mathbf{w}$  and  $b$ ?
- ▶ Need a quantity to optimise!
- ▶ Use the **margin**,  $\gamma$
- ▶ Maximise it!





Perpendicular distance from the decision boundary to the closest points on each side.

## Why maximise the margin?



- ▶ Maximum margin decision boundary (left) seems to better reflect the data characteristics than other boundary (right).

# Why maximise the margin?




- ▶ Maximum margin decision boundary (left) seems to better reflect the data characteristics than other boundary (right).
- ▶ Note how margin is much smaller on right and closest points have changed.
- ▶ There is going to be one 'best' boundary (w.r.t margin)
- ▶ **Statistical theory** justifying the choice.



## Computing the margin

$$\underline{2\gamma} = \frac{1}{\|\mathbf{w}\|} \mathbf{w}^T (\mathbf{x}_1 - \mathbf{x}_2)$$




## Computing the margin

$$2\gamma = \frac{1}{\|\mathbf{w}\|} \mathbf{w}^T (\mathbf{x}_1 - \mathbf{x}_2)$$

Fix the scale such that:

$$\mathbf{w}^T \mathbf{x}_1 + b = 1 \quad > 0$$

$$\mathbf{w}^T \mathbf{x}_2 + b = -1 \quad < 0$$



## Computing the margin

$$2\gamma = \frac{1}{\|\mathbf{w}\|} \boxed{\mathbf{w}^T(\mathbf{x}_1 - \mathbf{x}_2)}$$

$$= 2 \Rightarrow 2\gamma = \frac{1}{\|\mathbf{w}\|} \times 2$$

Fix the scale such that:

$$\mathbf{w}^T \mathbf{x}_1 + b = 1$$


$$\mathbf{w}^T \mathbf{x}_2 + b = -1$$

Therefore:

$$(\mathbf{w}^T \mathbf{x}_1 + b) - (\mathbf{w}^T \mathbf{x}_2 + b) = 2$$

$$+1 \quad \mathbf{w}^T(\mathbf{x}_1 - \mathbf{x}_2) = 2$$

$$\boxed{\gamma = \frac{1}{\|\mathbf{w}\|}}$$



# Maximising the margin

- ▶ We want to maximise  $\gamma = \frac{1}{\|\mathbf{w}\|}$

# Maximising the margin

- ▶ We want to maximise  $\gamma = \frac{1}{\|\mathbf{w}\|}$
- ▶ Equivalent to minimising  $\|\mathbf{w}\|$

# Maximising the margin

- ▶ We want to maximise  $\gamma = \frac{1}{\|\mathbf{w}\|}$
- ▶ Equivalent to minimising  $\|\mathbf{w}\|$
- ▶ Equivalent to minimising  $\frac{1}{2}\|\mathbf{w}\|^2 = \underline{\frac{1}{2}\mathbf{w}^T\mathbf{w}}$

# Maximising the margin

- ▶ We want to maximise  $\gamma = \frac{1}{\|\mathbf{w}\|}$
- ▶ Equivalent to minimising  $\|\mathbf{w}\|$
- ▶ Equivalent to minimising  $\frac{1}{2}\|\mathbf{w}\|^2 = \frac{1}{2}\mathbf{w}^T\mathbf{w}$
- ▶ There are some constraints:
  - ▶ For  $\mathbf{x}_n$  with  $t_n = 1$ :  $\mathbf{w}^T\mathbf{x}_n + b \geq 1$
  - ▶ For  $\mathbf{x}_n$  with  $t_n = -1$ :  $\mathbf{w}^T\mathbf{x}_n + b \leq -1$

# Maximising the margin

- ▶ We want to maximise  $\gamma = \frac{1}{\|\mathbf{w}\|}$
- ▶ Equivalent to minimising  $\|\mathbf{w}\|$
- ▶ Equivalent to minimising  $\frac{1}{2}\|\mathbf{w}\|^2 = \frac{1}{2}\mathbf{w}^T\mathbf{w}$
- ▶ There are some constraints:
  - ▶ For  $\mathbf{x}_n$  with  $t_n = 1$ :  $\mathbf{w}^T\mathbf{x}_n + b \geq 1$
  - ▶ For  $\mathbf{x}_n$  with  $t_n = -1$ :  $\mathbf{w}^T\mathbf{x}_n + b \leq -1$
- ▶ Which can be expressed more neatly as:

$t_n(\mathbf{w}^T\mathbf{x}_n + b) \geq 1$
- ▶ (This is why we use  $t_n = \pm 1$  and not  $t_n = \{0, 1\}$ .)



# Maximising the margin

- We have the following optimisation problem:

$$\operatorname{argmin}_{\mathbf{w}} \frac{1}{2} \mathbf{w}^T \mathbf{w}$$

$$\text{Subject to: } t_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1$$

# Maximising the margin

$$\frac{\partial (w^T w)}{\partial w} = 2w$$

- We have the following optimisation problem:

$$\underset{\mathbf{w}}{\operatorname{argmin}} \quad \frac{1}{2} \mathbf{w}^T \mathbf{w}$$

$$\text{Subject to: } t_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1$$

- Can put the constraints into the minimisation using **Lagrange multipliers**:

$$\underset{\mathbf{w}}{\operatorname{argmin}} \quad \frac{1}{2} \mathbf{w}^T \mathbf{w} - \sum_{n=1}^N \alpha_n (t_n(\mathbf{w}^T \mathbf{x}_n + b) - 1)$$

Subject to:  $\alpha_n \geq 0$

## What now?

- ▶ Let's think about what happens at the solution (we'll see why...)
- ▶ We know that  $\frac{\partial}{\partial \mathbf{w}} = 0$  and  $\frac{\partial}{\partial b} = 0$ .

## What now?

- ▶ Let's think about what happens at the solution (we'll see why...)
- ▶ We know that  $\frac{\partial}{\partial \mathbf{w}} = 0$  and  $\frac{\partial}{\partial b} = 0$ .

$$\frac{\partial}{\partial \mathbf{w}} = \mathbf{w} - \sum_n \alpha_n t_n \mathbf{x}_n = 0$$

$$\frac{\partial}{\partial b} = - \sum_n \alpha_n t_n = 0$$

- ▶ From which we can infer that:

$$\left. \begin{aligned} \mathbf{w} &= \sum_n \alpha_n t_n \mathbf{x}_n \\ \sum_n \alpha_n t_n &= 0 \end{aligned} \right\}$$

## What now?

- ▶ Let's think about what happens at the solution (we'll see why...)
- ▶ We know that  $\frac{\partial}{\partial \mathbf{w}} = 0$  and  $\frac{\partial}{\partial b} = 0$ .

$$\frac{\partial}{\partial \mathbf{w}} = \mathbf{w} - \sum_n \alpha_n t_n \mathbf{x}_n = 0$$

$$\frac{\partial}{\partial b} = - \sum_n \alpha_n t_n = 0$$

- ▶ From which we can infer that:

$$\mathbf{w} = \sum_n \alpha_n t_n \mathbf{x}_n$$

$$\sum_n \alpha_n t_n = 0$$

- ▶ Substitute these back into our optimisation problem:

$$\begin{aligned}
 & \frac{1}{2} \mathbf{w}^\top \mathbf{w} - \sum_n \alpha_n (t_n (\mathbf{w}^\top \mathbf{x}_n + b) - 1) \\
 & \quad \vdots \qquad \qquad \qquad \vdots \qquad \qquad \qquad \vdots \\
 = & \sum_n \alpha_n - \frac{1}{2} \sum_{n,m} \alpha_n \alpha_m t_n t_m \mathbf{x}_n^\top \mathbf{x}_m
 \end{aligned}$$

$\Rightarrow \dots + b \underbrace{\sum_n \alpha_n t_n}_0$

$$\frac{1}{2} \mathbf{w}^T \mathbf{w} - \sum_n \alpha_n (t_n (\mathbf{w}^T \mathbf{x}_n + b) - 1)$$

$$= \sum_n \alpha_n - \frac{1}{2} \sum_{n,m} \alpha_n \alpha_m t_n t_m \mathbf{x}_n^T \mathbf{x}_m$$

- ▶ Instead of minimising the previous expression, we can maximise this one (for reasons we won't go into).
- ▶ Subject to:

$$\alpha_n \geq 0$$

$$\sum_n \alpha_n t_n = 0$$

$$\begin{aligned}
& \frac{1}{2} \mathbf{w}^T \mathbf{w} - \sum_n \alpha_n (t_n (\mathbf{w}^T \mathbf{x}_n + b) - 1) \\
& \quad \vdots \qquad \qquad \qquad \vdots \qquad \qquad \qquad \vdots \\
& = \sum_n \alpha_n - \frac{1}{2} \sum_{n,m} \alpha_n \alpha_m t_n t_m \mathbf{x}_n^T \mathbf{x}_m
\end{aligned}$$


---

- ▶ Instead of minimising the previous expression, we can maximise this one (for reasons we won't go into).
- ▶ Subject to:

$$\left\{ \begin{array}{l} \alpha_n \geq 0 \\ \sum_n \alpha_n t_n = 0 \end{array} \right.$$

- ▶ Decision function was  $\text{sign}(\mathbf{w}^T \mathbf{x}_{\text{new}} + b)$  and is now:

$$t_{\text{new}} = \text{sign} \left( \sum_{n=1}^N \alpha_n t_n \mathbf{x}_n^T \mathbf{x}_{\text{new}} + b \right)$$

$w^T x_1 + b = +1$   
 $\Rightarrow b = 1 - w^T x_1$



So?

$$\begin{aligned} \operatorname{argmax}_{\alpha} \quad & \sum_{n=1}^N \alpha_n - \frac{1}{2} \sum_{n,m=1}^N \alpha_n \alpha_m t_n t_m \mathbf{x}_n^T \mathbf{x}_m \\ \text{subject to} \quad & \sum_{n=1}^N \alpha_n t_n = 0, \quad \alpha_n \geq 0 \end{aligned}$$



- ▶ This is a standard optimisation problem (quadratic programming)
- ▶ Has a single, global solution. This is very useful!
- ▶ Many algorithms around to solve it.
- ▶ e.g. quadprog in Matlab, CVXOPT in Python ...

So?

$$\begin{aligned} \operatorname{argmax}_{\alpha} \quad & \sum_{n=1}^N \alpha_n - \frac{1}{2} \sum_{n,m=1}^N \alpha_n \alpha_m t_n t_m \mathbf{x}_n^T \mathbf{x}_m \\ \text{subject to} \quad & \sum_{n=1}^N \alpha_n t_n = 0, \quad \alpha_n \geq 0 \end{aligned}$$

- ▶ This is a standard optimisation problem (quadratic programming)
- ▶ Has a single, global solution. This is very useful!
- ▶ Many algorithms around to solve it.
- ▶ e.g. quadprog in Matlab, CVXOPT in Python ...
- ▶ Once we have  $\alpha_n$ :

$$t_{\text{new}} = \operatorname{sign} \left( \sum_{n=1}^N \alpha_n t_n \mathbf{x}_n^T \mathbf{x}_{\text{new}} + b \right)$$

*Handwritten red note:*  $= \operatorname{sign}(w^T \mathbf{x}_{\text{new}} + b)$

# Primal and Dual

## Primal

$$\operatorname{argmin}_{\mathbf{w}} \frac{1}{2} \mathbf{w}^T \mathbf{w}$$

$$\text{Subject to: } t_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1$$

## Dual


$$\operatorname{argmax}_{\alpha} \sum_{n=1}^N \alpha_n - \frac{1}{2} \sum_{n,m=1}^N \alpha_n \alpha_m t_n t_m \mathbf{x}_n^T \mathbf{x}_m$$

$$\text{subject to } \sum_{n=1}^N \alpha_n t_n = 0, \quad \alpha_n \geq 0$$



- ▶ This is a standard optimisation problem (quadratic programming)
- ▶ Has a single, global solution. This is very useful!
- ▶ Many algorithms around to solve it.
- ▶ e.g. quadprog in Matlab, CVXOPT in Python ...


# Optimal boundary



- ▶ Optimisation gives us  $\alpha_1, \dots, \alpha_N$
- ▶ Compute  $\mathbf{w} = \sum_n \alpha_n t_n \mathbf{x}_n$
- ▶ Compute  $b = t_n - \mathbf{w}^T \mathbf{x}_n$  (for one of the closest points)
  - ▶ Recall that we defined  $\mathbf{w}^T \mathbf{x}_n + b = \pm 1 = t_n$  for closest points.
- ▶ Plot  $\mathbf{w}^T \mathbf{x} + b = 0$

# Support Vectors


- ▶ At the optimum, only 3 non-zero  $\alpha$  values (squares).



- ▶  $t_{\text{new}} = \text{sign} \left( \sum_n \alpha_n t_n \mathbf{x}_n^T \mathbf{x}_{\text{new}} + b \right)$
- ▶ Predictions only depend on these data-points!

# Support Vectors


- ▶ At the optimum, only 3 non-zero  $\alpha$  values (squares).



- ▶  $t_{\text{new}} = \text{sign} \left( \sum_n \alpha_n t_n \mathbf{x}_n^T \mathbf{x}_{\text{new}} + b \right)$
- ▶ Predictions only depend on these data-points!
- ▶ We knew that – margin is only a function of closest points.

# Support Vectors

- ▶ At the optimum, only 3 non-zero  $\alpha$  values (squares).




support vector

- ▶  $t_{\text{new}} = \text{sign} \left( \sum_n \alpha_n t_n \mathbf{x}_n^T \mathbf{x}_{\text{new}} + b \right) = \text{sign} \left( \mathbf{w}^T \mathbf{x}_{\text{new}} + b \right)$
- ▶ Predictions only depend on these data-points!
- ▶ We knew that – margin is only a function of closest points.
- ▶ These are called **Support Vectors**

# Support Vectors

- ▶ At the optimum, only 3 non-zero  $\alpha$  values (squares).




- ▶  $t_{\text{new}} = \text{sign} \left( \sum_n \alpha_n t_n \mathbf{x}_n^T \mathbf{x}_{\text{new}} + b \right)$
- ▶ Predictions only depend on these data-points!
- ▶ We knew that – margin is only a function of closest points.
- ▶ These are called **Support Vectors**
- ▶ Normally a small proportion of the data:
  - ▶ Solution is *sparse*.




# Is sparseness good?

► Not always:



# Is sparseness good?

- Not always:



- Why does this happen?

$$\underline{t_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1}$$

- All points must be on correct side of boundary.
- This is a *hard margin*

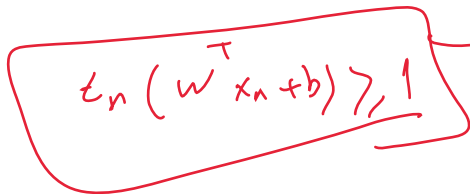
# Topics ...

- ▶ Linear SVM
- ▶ **Soft-Margin SVM**
- ▶ Kernels - Kernel SVM
- ▶ Classifier Performance

## Soft margin

- We can relax the constraints:

$$t_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1 - \xi_n, \quad \xi_n \geq 0$$


$$t_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1$$

## Soft margin

- We can relax the constraints:

$$t_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1 - \xi_n, \quad \xi_n \geq 0$$

- Our optimisation becomes:


$$\begin{aligned} & \underset{\mathbf{w}}{\operatorname{argmin}} \underbrace{\frac{1}{2} \mathbf{w}^T \mathbf{w}}_{\text{margin}} + C \underbrace{\sum_{n=1}^N \xi_n}_{\text{suppress constraints}} \\ & \text{subject to } \underbrace{t_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1 - \xi_n}_{\text{constraints}} \end{aligned}$$

## Soft margin

- We can relax the constraints:

$$t_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1 - \xi_n, \quad \xi_n \geq 0$$

- Our optimisation becomes:

$$\begin{aligned} & \underset{\mathbf{w}}{\operatorname{argmin}} \frac{1}{2} \mathbf{w}^T \mathbf{w} + C \sum_{n=1}^N \xi_n \\ & \text{subject to } t_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1 - \xi_n \end{aligned}$$


- And when we add Lagrange etc:

$$\begin{aligned} & \underset{\alpha}{\operatorname{argmax}} \sum_{n=1}^N \alpha_n - \frac{1}{2} \sum_{n,m=1}^N \alpha_n \alpha_m t_n t_m \mathbf{x}_n^T \mathbf{x}_m \\ & \text{subject to } \sum_{n=1}^N \alpha_n t_n = 0, \quad 0 \leq \alpha_n \leq C \end{aligned}$$

## Soft margin

- ▶ We can relax the constraints:

$$t_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1 - \xi_n, \quad \xi_n \geq 0$$

- ▶ Our optimisation becomes:

$$\begin{aligned} \underset{\mathbf{w}}{\operatorname{argmin}} \quad & \frac{1}{2} \mathbf{w}^T \mathbf{w} + C \sum_{n=1}^N \xi_n \\ \text{subject to} \quad & t_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1 - \xi_n \end{aligned}$$


- ▶ And when we add Lagrange etc:

$$\begin{aligned} \underset{\alpha}{\operatorname{argmax}} \quad & \sum_{n=1}^N \alpha_n - \frac{1}{2} \sum_{n,m=1}^N \alpha_n \alpha_m t_n t_m \mathbf{x}_n^T \mathbf{x}_m \\ \text{subject to} \quad & \sum_{n=1}^N \alpha_n t_n = 0, \quad 0 \leq \alpha_n \leq C \end{aligned}$$

- ▶ The **only** change is an upper-bound on  $\alpha_n$ !

## Soft margins

- Here's our problematic data again:




- $\alpha_n$  for the 'bad' square is 3.5.



## Soft margins

- ▶ Here's our problematic data again:




- ▶  $\alpha_n$  for the 'bad' square is 3.5.
- ▶ So, if we set  $C < 3.5$ , we should see this point having less influence and the boundary moving to somewhere more sensible...

## Soft margins

- Try  $C = 1$

$$0 \leq \alpha_n \leq C = 1$$



- We have an extra support vector.
- And a better decision boundary.

# Soft margins

- ▶ The choice of  $C$  is very important.
- ▶ Too high and we over-fit to noise.  $C$  is too large
- ▶ Too low and we underfit  
▶ ...and lose any sparsity.  $C$  is too small

# Soft margins

- ▶ The choice of  $C$  is very important.
- ▶ Too high and we *over-fit* to noise.
- ▶ Too low and we *underfit*
  - ▶ ...and lose any sparsity.
- ▶ Choose it using cross-validation.

# SVMs – some observations

- ▶ In our example, we started with 3 parameters:

$$\mathbf{w} = [w_1, w_2]^T, \quad b$$

- ▶ In general: D+1.

$$\mathbf{w}^T \mathbf{x} + b = 0$$

# SVMs – some observations

- ▶ In our example, we started with 3 parameters:

$$\mathbf{w} = [w_1, w_2]^T, \quad b$$

- ▶ In general:  $D+1$ .
- ▶ We now have  $N$ :  $\alpha_1, \dots, \alpha_N$

# SVMs – some observations

- ▶ In our example, we started with 3 parameters:

$$\mathbf{w} = [w_1, w_2]^T, \quad b$$

- ▶ In general:  $D+1$ .
- ▶ We now have  $N$ :  $\alpha_1, \dots, \alpha_N$
- ▶ Sounds harder?
- ▶ Depends on data dimensionality:
  - ▶ Typical Microarray dataset:
  - ▶  $D \sim 3000, N \sim 30$ .
  - ▶ In some cases  $N \ll D$

# Topics ...

- ▶ Linear SVM
- ▶ Soft-Margin SVM
- ▶ **Kernels - Kernel SVM**
- ▶ Classifier Performance



# Inner products

- ▶ Here's the optimisation problem:

$$\operatorname{argmax}_{\alpha} \sum_n \alpha_n - \frac{1}{2} \sum_{n,m} \alpha_n \alpha_m t_n t_m \underline{\mathbf{x}_n^T \mathbf{x}_m}$$

- ▶ Here's the decision function:

$$t_{\text{new}} = \operatorname{sign} \left( \sum_n \alpha_n t_n \underline{\mathbf{x}_n^T \mathbf{x}_{\text{new}}} + b \right)$$

# Inner products

- ▶ Here's the optimisation problem:

$$\operatorname{argmax}_{\alpha} \sum_n \alpha_n - \frac{1}{2} \sum_{n,m} \alpha_n \alpha_m t_n t_m \mathbf{x}_n^T \mathbf{x}_m$$

- ▶ Here's the decision function:


$$t_{\text{new}} = \operatorname{sign} \left( \sum_n \alpha_n t_n \mathbf{x}_n^T \mathbf{x}_{\text{new}} + b \right)$$

- ▶ Data ( $\mathbf{x}_n, \mathbf{x}_m, \mathbf{x}_{\text{new}}$ , etc) only appears as inner (dot) products:

$$\mathbf{x}_n^T \mathbf{x}_m, \mathbf{x}_n^T \mathbf{x}_{\text{new}}, \text{etc}$$


# Projections

- ▶ Our SVM can find linear decision boundaries.
- ▶ What if the data requires something nonlinear?



# Projections

- ▶ Our SVM can find linear decision boundaries.
- ▶ What if the data requires something nonlinear?



- ▶ We can transform the data e.g.:

$$\phi(\mathbf{x}_n) = x_{n1}^2 + x_{n2}^2$$

- ▶ So that it can be separated with a straight line.
- ▶ And use  $\phi(\mathbf{x}_n)$  instead of  $\mathbf{x}_n$  in our optimisation.

# Projections

- Our optimisation is now:

$$\operatorname{argmax}_{\alpha} \sum_n \alpha_n - \frac{1}{2} \sum_{n,m} \alpha_n \alpha_m t_n t_m \underbrace{\phi(\mathbf{x}_n)}^{\text{red}} \underbrace{\phi(\mathbf{x}_m)}^{\text{red}}$$

- And predictions:

$$t_{\text{new}} = \operatorname{sign} \left( \sum_n \alpha_n t_n \underbrace{\phi(\mathbf{x}_n)}^{\text{red}} \underbrace{\phi(\mathbf{x}_{\text{new}})}^{\text{red}} + b \right)$$

# Projections

- Our optimisation is now:

$$\operatorname{argmax}_{\alpha} \sum_n \alpha_n - \frac{1}{2} \sum_{n,m} \alpha_n \alpha_m t_n t_m \phi(\mathbf{x}_n)^T \phi(\mathbf{x}_m)$$

- And predictions:

$$t_{\text{new}} = \operatorname{sign} \left( \sum_n \alpha_n t_n \phi(\mathbf{x}_n)^T \phi(\mathbf{x}_{\text{new}}) + b \right)$$

- In this case:

$$\phi(\mathbf{x}_n)^T \phi(\mathbf{x}_m) = (x_{n1}^2 + x_{n2}^2)(x_{m1}^2 + x_{m2}^2) = k(\mathbf{x}_n, \mathbf{x}_m)$$

# Projections

- ▶ Our optimisation is now:

$$\operatorname{argmax}_{\alpha} \sum_n \alpha_n - \frac{1}{2} \sum_{n,m} \alpha_n \alpha_m t_n t_m \phi(\mathbf{x}_n)^T \phi(\mathbf{x}_m)$$

- ▶ And predictions:

$$t_{\text{new}} = \operatorname{sign} \left( \sum_n \alpha_n t_n \phi(\mathbf{x}_n)^T \phi(\mathbf{x}_{\text{new}}) + b \right)$$

- ▶ In this case:

$$\phi(\mathbf{x}_n)^T \phi(\mathbf{x}_m) = (x_{n1}^2 + x_{n2}^2)(x_{m1}^2 + x_{m2}^2) = k(\mathbf{x}_n, \mathbf{x}_m)$$

- ▶ We can think of the dot product in the projected space as a function of the original data.

# Projections

- ▶ We needn't directly think of projections at all. *def. of  $\phi(x_n)$*
- ▶ Can just think of functions  $k(\mathbf{x}_n, \mathbf{x}_m)$  that *are dot products in some space.*



# Projections

- ▶ We needn't directly think of projections at all.
- ▶ Can just think of functions  $k(\mathbf{x}_n, \mathbf{x}_m)$  that *are dot products in some space*.
- ▶ Called *kernel* functions.
- ▶ Don't ever need to actually project the data – just use the kernel function to compute what the dot product would be if we did project.

# Projections

- ▶ We needn't directly think of projections at all.
- ▶ Can just think of functions  $k(\mathbf{x}_n, \mathbf{x}_m)$  that *are dot products in some space*.
- ▶ Called *kernel* functions.
- ▶ Don't ever need to actually project the data – just use the kernel function to compute what the dot product would be if we did project.
- ▶ Optimisation task:

$$\operatorname{argmax}_{\alpha} \sum_n \alpha_n - \frac{1}{2} \sum_{n,m} \alpha_n \alpha_m t_n t_m \underline{k(\mathbf{x}_n, \mathbf{x}_m)}$$

# Projections

- ▶ We needn't directly think of projections at all.
- ▶ Can just think of functions  $k(\mathbf{x}_n, \mathbf{x}_m)$  that *are dot products in some space*.
- ▶ Called *kernel* functions.
- ▶ Don't ever need to actually project the data – just use the kernel function to compute what the dot product would be if we did project.
- ▶ Optimisation task:

$$\operatorname{argmax}_{\alpha} \sum_n \alpha_n - \frac{1}{2} \sum_{n,m} \alpha_n \alpha_m t_n t_m k(\mathbf{x}_n, \mathbf{x}_m)$$

- ▶ Predictions:

$$t_{\text{new}} = \operatorname{sign} \left( \sum_n \alpha_n t_n k(\mathbf{x}_n, \mathbf{x}_{\text{new}}) + b \right)$$

# Kernels

- ▶ Plenty of off-the-shelf kernels that we can use:

# Kernels

- ▶ Plenty of off-the-shelf kernels that we can use:
- ▶ Linear:

$$k(\mathbf{x}_n, \mathbf{x}_m) = \mathbf{x}_n^T \mathbf{x}_m$$

# Kernels

- ▶ Plenty of off-the-shelf kernels that we can use:
- ▶ Linear:

$$k(\mathbf{x}_n, \mathbf{x}_m) = \mathbf{x}_n^T \mathbf{x}_m$$

- ▶ Gaussian:

$$k(\mathbf{x}_n, \mathbf{x}_m) = \exp \left\{ -\beta (\mathbf{x}_n - \mathbf{x}_m)^T (\mathbf{x}_n - \mathbf{x}_m) \right\}$$

# Kernels

- ▶ Plenty of off-the-shelf kernels that we can use:
- ▶ Linear:

$$k(\mathbf{x}_n, \mathbf{x}_m) = \mathbf{x}_n^T \mathbf{x}_m$$

- ▶ Gaussian:

$$k(\mathbf{x}_n, \mathbf{x}_m) = \exp \left\{ -\beta (\mathbf{x}_n - \mathbf{x}_m)^T (\mathbf{x}_n - \mathbf{x}_m) \right\}$$

- ▶ Polynomial:

$$k(\mathbf{x}_n, \mathbf{x}_m) = (1 + \mathbf{x}_n^T \mathbf{x}_m)^\beta$$

# Kernels

- ▶ Plenty of off-the-shelf kernels that we can use:
- ▶ Linear:

$$k(\mathbf{x}_n, \mathbf{x}_m) = \mathbf{x}_n^T \mathbf{x}_m$$

- ▶ Gaussian:

$$k(\mathbf{x}_n, \mathbf{x}_m) = \exp \left\{ -\beta (\mathbf{x}_n - \mathbf{x}_m)^T (\mathbf{x}_n - \mathbf{x}_m) \right\}$$

- ▶ Polynomial:

$$k(\mathbf{x}_n, \mathbf{x}_m) = (1 + \mathbf{x}_n^T \mathbf{x}_m)^\beta$$

- ▶ These all correspond to  $\phi(\mathbf{x}_n)^T \phi(\mathbf{x}_m)$  for some transformation  $\phi(\mathbf{x}_n)$ .
- ▶ Don't know what the projections  $\phi(\mathbf{x}_n)$  are – don't need to know!



# Kernels

- ▶ Our algorithm is still only finding linear boundaries....

# Kernels

- ▶ Our algorithm is still only finding linear boundaries....
- ▶ ...but we're finding linear boundaries in some other space.

# Kernels

- ▶ Our algorithm is still only finding linear boundaries....
- ▶ ...but we're finding linear boundaries in some other space.
- ▶ The optimisation is just as simple, regardless of the kernel choice.
  - ▶ Still a quadratic program.
  - ▶ Still a single, global optimum.



# Kernels

- ▶ Our algorithm is still only finding linear boundaries....
- ▶ ...but we're finding linear boundaries in some other space.
- ▶ The optimisation is just as simple, regardless of the kernel choice.
  - ▶ Still a quadratic program.
  - ▶ Still a single, global optimum.
- ▶ We can find very complex decision boundaries with a linear algorithm!

## A technical point

- ▶ Our decision boundary was defined as  $\mathbf{w}^T \mathbf{x} + b = 0$ .
- ▶ Now,  $\mathbf{w}$  is defined as:

$$\rightarrow \mathbf{w} = \sum_{n=1}^N \alpha_n t_n \phi(\mathbf{x}_n)$$

- ▶ We don't know  $\phi(\mathbf{x}_n)$ .

## A technical point

- ▶ Our decision boundary was defined as  $\mathbf{w}^T \mathbf{x} + b = 0$ .
- ▶ Now,  $\mathbf{w}$  is defined as:

$$\mathbf{w} = \sum_{n=1}^N \alpha_n t_n \phi(\mathbf{x}_n)$$

- ▶ We don't know  $\phi(\mathbf{x}_n)$ .
- ▶ We **only know**  $\phi(\mathbf{x}_n)^T \phi(\mathbf{x}_m) = k(\mathbf{x}_n, \mathbf{x}_m)$
- ▶ So, we can't compute  $\mathbf{w}$  or the boundary!

## A technical point

- ▶ Our decision boundary was defined as  $\mathbf{w}^T \mathbf{x} + b = 0$ .
- ▶ Now,  $\mathbf{w}$  is defined as:

$$\mathbf{w} = \sum_{n=1}^N \alpha_n t_n \phi(\mathbf{x}_n)$$

- ▶ We don't know  $\phi(\mathbf{x}_n)$ .
- ▶ We **only know**  $\phi(\mathbf{x}_n)^T \phi(\mathbf{x}_m) = k(\mathbf{x}_n, \mathbf{x}_m)$
- ▶ So, we can't compute  $\mathbf{w}$  or the boundary!
- ▶ But we can evaluate the predictions on a grid of  $\mathbf{x}_{\text{new}}$  and use Python/Matlab to draw a contour:

$$\sum_{n=1}^N \alpha_n t_n k(\mathbf{x}_n, \mathbf{x}_{\text{new}}) + b$$



## Aside: kernelising other algorithms

- ▶ **Many** algorithms can be kernelised.
  - ▶ Any that can be written with data only appearing as inner products.
- ▶ Simple algorithms can be used to solve very complex problems!
- ▶ Class exercise:
  - ▶ KNN requires the distance between  $\mathbf{x}_{\text{new}}$  and each  $\mathbf{x}_n$ :

$$(\mathbf{x}_{\text{new}} - \mathbf{x}_n)^T (\mathbf{x}_{\text{new}} - \mathbf{x}_n)$$

- ▶ Can we kernelise it?


$$= \mathbf{x}_{\text{new}}^T \mathbf{x}_{\text{new}} - 2 \mathbf{x}_n^T \mathbf{x}_{\text{new}} + \mathbf{x}_n^T \mathbf{x}_n$$

Handwritten red annotations below the equation:

- An arrow points from  $\mathbf{x}_{\text{new}}^T \mathbf{x}_{\text{new}}$  down to  $k(\cdot, \cdot)$ .
- An arrow points from  $\mathbf{x}_n^T \mathbf{x}_{\text{new}}$  down to  $k(\cdot, \cdot)$ .
- An arrow points from  $\mathbf{x}_n^T \mathbf{x}_n$  down to  $k(\cdot, \cdot)$ .



## Example – nonlinear data




- ▶ We'll use a Gaussian kernel:

$$k(\mathbf{x}_n, \mathbf{x}_m) = \exp \left\{ -\beta (\mathbf{x}_n - \mathbf{x}_m)^T (\mathbf{x}_n - \mathbf{x}_m) \right\}$$

A red arrow points to the parameter  $\beta$  in the equation.

- ▶ And vary  $\beta$  ( $C = 10$ ).


## Examples



►  $\beta = 1$ .

$$k(\mathbf{x}_n, \mathbf{x}_m) = \exp \left\{ -\beta (\mathbf{x}_n - \mathbf{x}_m)^\top (\mathbf{x}_n - \mathbf{x}_m) \right\}$$


## Examples



►  $\beta = 0.01$ .

$$k(\mathbf{x}_n, \mathbf{x}_m) = \exp \left\{ -\beta (\mathbf{x}_n - \mathbf{x}_m)^\top (\mathbf{x}_n - \mathbf{x}_m) \right\}$$

## Examples



►  $\beta = 50$ .

$$k(\mathbf{x}_n, \mathbf{x}_m) = \exp \left\{ -\beta (\mathbf{x}_n - \mathbf{x}_m)^T (\mathbf{x}_n - \mathbf{x}_m) \right\}$$

# The Gaussian kernel

- ▶  $\beta$  controls the *complexity* of the decision boundaries.

# The Gaussian kernel

- ▶  $\beta$  controls the *complexity* of the decision boundaries.
- ▶  $\beta = 0.01$  was too simple:
  - ▶ Not flexible enough to surround just the square class.

# The Gaussian kernel


- ▶  $\beta$  controls the *complexity* of the decision boundaries.
- ▶  $\beta = 0.01$  was too simple:
  - ▶ Not flexible enough to surround just the square class.
- ▶  $\beta = 50$  was too complex:
  - ▶ *Memorises* the data.

# The Gaussian kernel

- ▶  $\beta$  controls the *complexity* of the decision boundaries.
- ▶  $\beta = 0.01$  was too simple:
  - ▶ Not flexible enough to surround just the square class.
- ▶  $\beta = 50$  was too complex:
  - ▶ *Memorises* the data.
- ▶  $\beta = 1$  was about right.



# The Gaussian kernel

- ▶  $\beta$  controls the *complexity* of the decision boundaries.
  - ▶  $\beta = 0.01$  was too simple:
    - ▶ Not flexible enough to surround just the square class.
  - ▶  $\beta = 50$  was too complex:
    - ▶ *Memorises* the data.
  - ▶  $\beta = 1$  was about right.
  - ▶ Neither  $\beta = 50$  or  $\beta = 0.01$  will *generalise* well.
  - ▶ Both are also non-sparse (lots of support vectors).
- 

# Choosing kernel function, parameters and $C$

- ▶ Kernel function and parameter choice is data dependent.
- ▶ Easy to overfit.

# Choosing kernel function, parameters and $C$

- ▶ Kernel function and parameter choice is data dependent.
- ▶ Easy to overfit.
- ▶ Need to set  $C$  too
- ▶  $C$  and  $\beta$  are *linked*
  - ▶  $C$  too high – overfitting.
  - ▶  $C$  too low – underfitting.

# Choosing kernel function, parameters and $C$

- ▶ Kernel function and parameter choice is data dependent.
- ▶ Easy to overfit.
- ▶ Need to set  $C$  too
- ▶  $C$  and  $\beta$  are *linked*
  - ▶  $C$  too high – overfitting.
  - ▶  $C$  too low – underfitting.
- ▶ Cross-validation!

# Choosing kernel function, parameters and $C$

- ▶ Kernel function and parameter choice is data dependent.
- ▶ Easy to overfit.
- ▶ Need to set  $C$  too
- ▶  $C$  and  $\beta$  are *linked*
  - ▶  $C$  too high – overfitting.
  - ▶  $C$  too low – underfitting.
- ▶ Cross-validation!
- ▶ Search over  $\beta$  and  $C$ 
  - ▶ SVM scales with  $N^3$  (naive implementation)
  - ▶ For large  $N$ , cross-validation over many  $C$  and  $\beta$  values is infeasible.

# Summary - SVMs

- ▶ Described a classifier that is optimised by maximising the *margin*.
- ▶ Did some re-arranging to turn it into a quadratic programming problem.
- ▶ Loosened the SVM constraints to allow points on the wrong side of boundary.
- ▶ Saw that data only appear as inner products.
- ▶ Introduced the idea of kernels.
- ▶ Can fit a linear boundary in some other space without explicitly projecting.
- ▶ Other algorithms can be kernelised...we'll see a clustering one in the future.

# Topics ...

- ▶ Linear SVM
- ▶ Soft-Margin SVM
- ▶ Kernels - Kernel SVM
- ▶ **Classifier Performance**

# Performance evaluation

- ▶ We've seen 3 classification algorithms, more will come later ...
- ▶ How do we choose?
  - ▶ Which algorithm?
  - ▶ Which parameters?
- ▶ Need performance indicators.



# Performance evaluation

- ▶ We've seen 3 classification algorithms, more will come later ...
- ▶ How do we choose?
  - ▶ Which algorithm?
  - ▶ Which parameters?
- ▶ Need performance indicators.
- ▶ We'll cover:
  - ▶ 0/1 loss.
  - ▶ ROC analysis (sensitivity and specificity)
  - ▶ Confusion matrices

## 0/1 loss

- ▶ 0/1 loss: proportion of times classifier is wrong.
- ▶ Consider a set of predictions  $t_1, \dots, t_N$  and a set of true labels  $t_1^*, \dots, t_N^*$ .
- ▶ Mean loss is defined as:

$$\frac{1}{N} \sum_{n=1}^N \delta(t_n \neq t_n^*)$$

- ▶ ( $\delta(a)$  is 1 if  $a$  is true and 0 otherwise)

$\delta(\text{cond}) = 1$   
if cond is True  
0 otherwise.

## 0/1 loss

- ▶ 0/1 loss: proportion of times classifier is wrong.
- ▶ Consider a set of predictions  $t_1, \dots, t_N$  and a set of true labels  $t_1^*, \dots, t_N^*$ .
- ▶ Mean loss is defined as:

$$\frac{1}{N} \sum_{n=1}^N \delta(t_n \neq t_n^*)$$

- ▶ ( $\delta(a)$  is 1 if  $a$  is true and 0 otherwise)
- ▶ Advantages:
  - ▶ Can do binary or multiclass classification.
  - ▶ Simple to compute.
  - ▶ Single value.

## 0/1 loss

Disadvantage: Doesn't take into account class imbalance:

## 0/1 loss

Disadvantage: Doesn't take into account class imbalance:

- ▶ We're building a classifier to detect a rare disease.
- ▶ Assume only 1% of population is diseased.

## 0/1 loss

Disadvantage: Doesn't take into account class imbalance:

- ▶ We're building a classifier to detect a rare disease.
- ▶ Assume only 1% of population is diseased.
- ▶ Diseased:  $t = 1$
- ▶ Healthy:  $t = 0$

## 0/1 loss

Disadvantage: Doesn't take into account class imbalance:

- ▶ We're building a classifier to detect a rare disease.
- ▶ Assume only 1% of population is diseased.
- ▶ Diseased:  $t = 1$  1%  $\implies 10/1000 = 0.01$
- ▶ Healthy:  $t = 0$  99%
- ▶ What if we always predict healthy? ( $t = 0$ )

## 0/1 loss

Disadvantage: Doesn't take into account class imbalance:

- ▶ We're building a classifier to detect a rare disease.
- ▶ Assume only 1% of population is diseased.
- ▶ Diseased:  $t = 1$
- ▶ Healthy:  $t = 0$
- ▶ What if we always predict healthy? ( $t = 0$ )
- ▶ Accuracy 99%
- ▶ But classifier is rubbish!



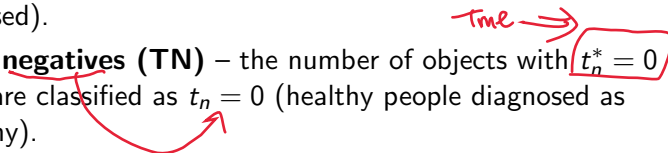
# Sensitivity and specificity

- ▶ We'll stick with our disease example.
- ▶ Need to define 4 quantities. The numbers of:

# Sensitivity and specificity

- ▶ We'll stick with our disease example.
- ▶ Need to define 4 quantities. The numbers of:
- ▶ **True positives (TP)** – the number of objects with  $t_n^* = 1$  that are classified as  $t_n = 1$  (diseased people diagnosed as diseased).

# Sensitivity and specificity

- ▶ We'll stick with our disease example.
  - ▶ Need to define 4 quantities. The numbers of:
  - ▶ True positives (TP) – the number of objects with  $t_n^* = 1$  that are classified as  $t_n = 1$  (diseased people diagnosed as diseased).
  - ▶ **True negatives (TN)** – the number of objects with  $t_n^* = 0$  that are classified as  $t_n = 0$  (healthy people diagnosed as healthy).
- 

# Sensitivity and specificity

- ▶ We'll stick with our disease example.
- ▶ Need to define 4 quantities. The numbers of:
- ▶ True positives (TP) – the number of objects with  $t_n^* = 1$  that are classified as  $t_n = 1$  (diseased people diagnosed as diseased).
- ▶ True negatives (TN) – the number of objects with  $t_n^* = 0$  that are classified as  $t_n = 0$  (healthy people diagnosed as healthy).
- ▶ **False positives (FP)** – the number of objects with  $t_n^* = 0$  that are classified as  $t_n = 1$  (healthy people diagnosed as diseased).

# Sensitivity and specificity

- ▶ We'll stick with our disease example.
- ▶ Need to define 4 quantities. The numbers of:
- ▶ True positives (TP) – the number of objects with  $t_n^* = 1$  that are classified as  $t_n = 1$  (diseased people diagnosed as diseased).
- ▶ True negatives (TN) – the number of objects with  $t_n^* = 0$  that are classified as  $t_n = 0$  (healthy people diagnosed as healthy).
- ▶ False positives (FP) – the number of objects with  $t_n^* = 0$  that are classified as  $t_n = 1$  (healthy people diagnosed as diseased).
- ▶ **False negatives (FN)** – the number of objects with  $t_n^* = 1$  that are classified as  $t_n = 0$  (diseased people diagnosed as healthy).

# Sensitivity

$$S_e = \frac{TP}{TP + FN}$$

Handwritten annotations in red:

- Arrow from  $TP$  to  $\{t_p^* = 1, t_n = 1\}$
- Arrow from  $TP + FN$  to  $\{t_p^* = 1, t_n = 0\}$
- Arrow from  $TP + FN$  to  $\{t_n^* = 1, t_n = 1\}$

- ▶ The proportion of diseased people that we classify as diseased.
- ▶ The higher the better.
- ▶ In our example,  $S_e = 0$ .

Handwritten annotation in red:

$$\{t_n^* = 1\}$$

# Specificity

$$S_p = \frac{TN}{TN + FP} = 1$$

Handwritten annotations in red:

- $\{t_n^* = 0, t_n = 0\}$  (left)
- $\{t_n^* = 0, t_n = 0\}$  (top right)
- $\{t_n^* = 0, t_n = 1\}$  (bottom right)
- $\{t_n^* = 0\}$  (bottom center)

Arrows indicate the mapping from the sets to the terms in the equation:  $\{t_n^* = 0, t_n = 0\}$  to  $TN$ ,  $\{t_n^* = 0, t_n = 1\}$  to  $FP$ , and  $\{t_n^* = 0\}$  to the denominator  $TN + FP$ .

- ▶ The proportion of healthy people that we classify as healthy.
- ▶ The higher the better.
- ▶ In our example,  $S_p = 1$ .

# Optimising sensitivity and specificity

- ▶ We would like both to be as high as possible.
- ▶ Often increasing one will decrease the other.



# Optimising sensitivity and specificity

- ▶ We would like both to be as high as possible.
- ▶ Often increasing one will decrease the other.
- ▶ Balance will depend on application:
- ▶ e.g. diagnosis:
  - ▶ We can probably tolerate a decrease in specificity (healthy people diagnosed as diseased)....
  - ▶ ...if it gives us an increase in sensitivity (getting diseased people right).

# ROC analysis

other  
thresholds.

- ▶ Many classification algorithms involve setting a threshold.
- ▶ e.g. SVM:

$$t_{\text{new}} = \text{sign} \left( \sum_{n=1}^N t_n \alpha_n k(\mathbf{x}_n, \mathbf{x}_{\text{new}}) + b \right)$$

- ▶ Implies a threshold of zero (sign function)




# ROC analysis

- ▶ Many classification algorithms involve setting a threshold.
- ▶ e.g. SVM:


$$t_{\text{new}} = \text{sign} \left( \sum_{n=1}^N t_n \alpha_n k(\mathbf{x}_n, \mathbf{x}_{\text{new}}) + b \right)$$

- ▶ Implies a threshold of zero (sign function)
- ▶ However, we could use any threshold we like....
- ▶ The **Receiver Operating Characteristic (ROC) curve** shows how  $S_e$  and  $1 - S_p$  vary as the threshold changes.




- ▶ SVM for nonlinear data with  $\beta = 50$ .
- ▶ Each point is a threshold value.
- ▶ Bottom left – everything classified as 0 (-1 in SVM)
- ▶ Top right – everything classified as 1.
- ▶ Goal: get the curve to the top left corner – perfect classification ( $S_e = 1, S_p = 1$ ).

## ROC curve



- ▶ SVM for nonlinear data with  $\beta = 0.01$ .
- ▶ Better than  $\beta = 50$ 
  - ▶ Closer to top left corner.

## ROC curve



- ▶ SVM for nonlinear data with  $\beta = 1$ .
- ▶ Better still.

# AUC

- ▶ We can quantify performance by computing the **Area Under the ROC Curve (AUC)**
- ▶ The higher this value, the better.
  - ▶  $\beta = 50$ : AUC=0.8348
  - ▶  $\beta = 0.01$ : AUC= 0.9551
  - ▶  $\beta = 1$ : AUC=0.9936

# AUC

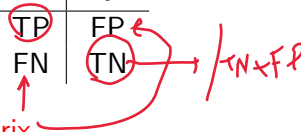
- ▶ We can quantify performance by computing the **Area Under the ROC Curve (AUC)**
- ▶ The higher this value, the better.
  - ▶  $\beta = 50$ : AUC=0.8348
  - ▶  $\beta = 0.01$ : AUC= 0.9551
  - ▶  $\beta = 1$ : AUC=0.9936
- ▶ AUC is generally a safer measure than 0/1 loss.



# Confusion matrices

The quantities we used to compute  $S_e$  and  $S_p$  can be neatly summarised in a table:

		True class	
		1	0
Predicted class	1	TP	FP
	0	FN	TN



- ▶ This is known as a **confusion matrix**
- ▶ It is particularly useful for multi-class classification.
- ▶ Tells us where the mistakes are being made.
- ▶ Note that normalising columns gives us  $S_e$  and  $S_p$

# Confusion matrices – example

- ▶ 20 newsgroups data.
- ▶ Thousands of documents from 20 classes (newsgroups)
- ▶ Use a Naive Bayes classifier ( $\approx 50000$  dimensions (words)!)
  - ▶ Details in book Chapter.
- ▶  $\approx 7000$  independent test documents.
- ▶ Summarise results in  $20 \times 20$  confusion matrix:

			True class										
			10	11	12	13	14	15	16	17	18	19	20
Predicted class	1	...	4	2	0	2	10	4	7	1	12	7	47
	2	...	0	0	4	18	7	8	2	0	1	1	3
	3	...	0	0	1	0	1	0	1	0	0	0	0
	4	...	1	0	1	28	3	0	0	0	0	0	0
	⋮												
	16	...	3	2	2	5	17	4	376	3	7	2	68
	17	...	1	0	9	0	3	1	3	325	3	95	19
	18	...	2	1	0	2	6	2	1	2	325	4	5
	19	...	8	4	8	0	10	21	1	16	19	185	7
	20	...	0	0	1	0	1	1	2	4	0	1	92

		True class											
		...	10	11	12	13	14	15	16	17	18	19	20
Predicted class	1	...	4	2	0	2	10	4	7	1	12	7	47
	2	...	0	0	4	18	7	8	2	0	1	1	3
	3	...	0	0	1	0	1	0	1	0	0	0	0
	4	...	1	0	1	28	3	0	0	0	0	0	0
	...												
	16	...	3	2	2	5	17	4	376	3	7	2	<b>68</b>
	17	...	1	0	9	0	3	1	3	325	3	<b>95</b>	19
	18	...	2	1	0	2	6	2	1	2	325	4	5
	19	...	8	4	8	0	10	21	1	16	19	<b>185</b>	7
	20	...	0	0	1	0	1	1	2	4	0	1	<b>92</b>

- ▶ Algorithm is getting 'confused' between classes 20 and 16, and 19 and 17.
  - ▶ 17: talk.politics.guns
  - ▶ 19: talk.politics.misc

		True class											
		...	10	11	12	13	14	15	16	17	18	19	20
Predicted class	1	...	4	2	0	2	10	4	7	1	12	7	47
	2	...	0	0	4	18	7	8	2	0	1	1	3
	3	...	0	0	1	0	1	0	1	0	0	0	0
	4	...	1	0	1	28	3	0	0	0	0	0	0
	⋮												
	16	...	3	2	2	5	17	4	376	3	7	2	<b>68</b>
	17	...	1	0	9	0	3	1	3	325	3	<b>95</b>	19
	18	...	2	1	0	2	6	2	1	2	325	4	5
	19	...	8	4	8	0	10	21	1	16	19	<b>185</b>	7
	20	...	0	0	1	0	1	1	2	4	0	1	<b>92</b>

► Algorithm is getting 'confused' between classes 20 and 16, and 19 and 17.

- 17: talk.politics.guns
- 19: talk.politics.misc
- 16: talk.religion.misc
- 20: soc.religion.christian

			True class										
			10	11	12	13	14	15	16	17	18	19	20
Predicted class	1	...	4	2	0	2	10	4	7	1	12	7	47
	2	...	0	0	4	18	7	8	2	0	1	1	3
	3	...	0	0	1	0	1	0	1	0	0	0	0
	4	...	1	0	1	28	3	0	0	0	0	0	0
	⋮												
	16	...	3	2	2	5	17	4	376	3	7	2	68
	17	...	1	0	9	0	3	1	3	325	3	95	19
	18	...	2	1	0	2	6	2	1	2	325	4	5
	19	...	8	4	8	0	10	21	1	16	19	185	7
	20	...	0	0	1	0	1	1	2	4	0	1	92

- ▶ Algorithm is getting 'confused' between classes 20 and 16, and 19 and 17.
  - ▶ 17: talk.politics.guns
  - ▶ 19: talk.politics.misc
  - ▶ 16: talk.religion.misc
  - ▶ 20: soc.religion.christian
- ▶ Maybe these should be just one class?
- ▶ Maybe we need more data in these classes?

		True class											
		...	10	11	12	13	14	15	16	17	18	19	20
Predicted class	1	...	4	2	0	2	10	4	7	1	12	7	47
	2	...	0	0	4	18	7	8	2	0	1	1	3
	3	...	0	0	1	0	1	0	1	0	0	0	0
	4	...	1	0	1	28	3	0	0	0	0	0	0
		...											
Predicted class	16	...	3	2	2	5	17	4	376	3	7	2	68
	17	...	1	0	9	0	3	1	3	325	3	95	19
	18	...	2	1	0	2	6	2	1	2	325	4	5
	19	...	8	4	8	0	10	21	1	16	19	185	7
	20	...	0	0	1	0	1	1	2	4	0	1	92

- ▶ Algorithm is getting 'confused' between classes 20 and 16, and 19 and 17.
  - ▶ 17: talk.politics.guns
  - ▶ 19: talk.politics.misc
  - ▶ 16: talk.religion.misc
  - ▶ 20: soc.religion.christian
- ▶ Maybe these should be just one class?
- ▶ Maybe we need more data in these classes?
- ▶ Confusion matrix helps us direct our efforts to improving the classifier.

# Summary

- ▶ SVM: a kernel classifier.
- ▶ Linear classifier – (possibly) nonlinear data transformation.



# Summary

- ▶ SVM: a kernel classifier.
- ▶ Linear classifier – (possibly) nonlinear data transformation.
- ▶ Introduced two different performance measures:
  - ▶ 0/1 loss
  - ▶ ROC/AUC

# Summary

- ▶ SVM: a kernel classifier.
- ▶ Linear classifier – (possibly) nonlinear data transformation.
- ▶ Introduced two different performance measures:
  - ▶ 0/1 loss
  - ▶ ROC/AUC
- ▶ Introduced confusion matrices – a way of assessing the performance of a multi-class classifier.